

# Getting started

to start use pandas libaray

```
In [2]: import pandas as pd
```

## DataFrame

*A DataFrame is a table. It contains an array of individual entries, each of which has a certain value. Each entry corresponds to a row (or record) and a column.*

We are using the `pd.DataFrame()` constructor to generate these DataFrame objects. The syntax for declaring a new one is a dictionary whose keys are the column names, and whose values are a list of entries. This is the standard way of constructing a new DataFrame, and the one you are most likely to encounter.

```
In [4]: pd.DataFrame({'Yes': [50, 21], 'No': [131, 2]})
```

```
Out[4]:
```

	Yes	No
0	50	131
1	21	2

```
In [5]: pd.DataFrame({'Bob': ['I liked it.', 'It was awful.'], 'Sue': ['Pretty good.', 'Bland.]
```

```
Out[5]:
```

	Bob	Sue
0	I liked it.	Pretty good.
1	It was awful.	Bland.

The list of row labels used in a DataFrame is known as an Index. We can assign values to it by using an index parameter in our constructor:

```
In [6]: pd.DataFrame({'Bob': ['I liked it.', 'It was awful.'],
                      'Sue': ['Pretty good.', 'Bland.'],
                      index=['Product A', 'Product B'])
```

```
Out[6]:
```

	Bob	Sue
Product A	I liked it.	Pretty good.
Product B	It was awful.	Bland.

# Series

is a sequence of data values, is a list. And in fact you can create one with nothing more than a list:

```
In [7]: pd.Series([1, 2, 3, 4, 5])
```

```
Out[7]: 0    1
        1    2
        2    3
        3    4
        4    5
        dtype: int64
```

Series is, in essence, a single column of a DataFrame. So you can assign row labels to the Series the same way as before, using an index parameter. However, a Series does not have a column name, it only has one overall name:

```
In [8]: pd.Series([30, 35, 40], index=['2015 Sales', '2016 Sales', '2017 Sales'], name='Product A')
```

```
Out[8]: 2015 Sales    30
        2016 Sales    35
        2017 Sales    40
        Name: Product A, dtype: int64
```

---

## Reading data files

Data can be stored in any of a number of different forms and formats. By far the most basic of these is the humble CSV file. When you open a CSV file you get something that looks like this:

```
In [3]: reviews = pd.read_csv('win-data.csv')
```

```
In [4]: reviews
```

```
Out[4]:
```

	Unnamed: 0	country	description	designation	points	price	province	region_1	region_2
0	0	US	This tremendous 100% varietal wine hails from ...	Martha's Vineyard	96	235.0	California	Napa Valley	Napa

	Unnamed: 0	country	description	designation	points	price	province	region_1	region
1	1	Spain	Ripe aromas of fig, blackberry and cassis are ...	Carodorum Selección Especial Reserva	96	110.0	Northern Spain	Toro	Na
2	2	US	Mac Watson honors the memory of a wine once ma...	Special Selected Late Harvest	96	90.0	California	Knights Valley	Sonor
3	3	US	This spent 20 months in 30% new French oak, an...	Reserve	96	65.0	Oregon	Willamette Valley	Willamet Vall
4	4	France	This is the top wine from La Bégude, named aft...	La Brûlade	95	66.0	Provence	Bandol	Na
...	...	...	...	...	...	...	...	...	...
150925	150925	Italy	Many people feel Fiano represents southern Ita...	NaN	91	20.0	Southern Italy	Fiano di Avellino	Na
150926	150926	France	Offers an intriguing nose with ginger, lime an...	Cuvée Prestige	91	27.0	Champagne	Champagne	Na
150927	150927	Italy	This classic example comes from a cru vineyard...	Terre di Dora	91	20.0	Southern Italy	Fiano di Avellino	Na
150928	150928	France	A perfect salmon shade, with scents of peaches...	Grand Brut Rosé	90	52.0	Champagne	Champagne	Na
150929	150929	Italy	More Pinot Grigios should taste like this. A r...	NaN	90	15.0	Northeastern Italy	Alto Adige	Na

150930 rows × 11 columns



In [22]:

```
data.head()
```

Out[22]:

	Unnamed: 0	country	description	designation	points	price	province	region_1	region_2	vari
0	0	US	This tremendous 100% varietal wine hails from ...	Martha's Vineyard	96	235.0	California	Napa Valley	Napa	Cabel Sauvig
1	1	Spain	Ripe aromas of fig, blackberry and cassis are ...	Carodorum Selección Especial Reserva	96	110.0	Northern Spain	Toro	NaN	Tinta 1
2	2	US	Mac Watson honors the memory of a wine once ma...	Special Selected Late Harvest	96	90.0	California	Knights Valley	Sonoma	Sauvig Bl
3	3	US	This spent 20 months in 30% new French oak, an...	Reserve	96	65.0	Oregon	Willamette Valley	Willamette Valley	Pinot f
4	4	France	This is the top wine from La Bégude, named aft...	La Brûlade	95	66.0	Provence	Bandol	NaN	Prove red bl



## Indexing, Selecting & Assigning

Selecting specific values of a pandas DataFrame or Series to work on is an implicit step in almost any data operation.

Hence to access the country property of data we can use:

In [23]:

```
reviews.country
```

Out[23]:

```
0      US
1    Spain
```

```

2          US
3          US
4      France
...
150925    Italy
150926    France
150927    Italy
150928    France
150929    Italy
Name: country, Length: 150930, dtype: object

```

**for columns in a DataFrame: we can access its values using the indexing ([ ]) operator.**

```
In [24]: reviews['country']
```

```

Out[24]: 0          US
1      Spain
2          US
3          US
4      France
...
150925    Italy
150926    France
150927    Italy
150928    France
150929    Italy
Name: country, Length: 150930, dtype: object

```

**to drill down to a single specific value, we need only use the indexing operator [ ] once more:**

```
In [26]: reviews['country'][0]
```

```
Out[26]: 'US'
```

## Indexing

### Index-based selection

**Pandas indexing works in one of two paradigms. The first is index-based selection: selecting data based on its numerical position in the data. `iloc` follows this paradigm.**

```
In [27]: reviews.iloc[0]
```

```

Out[27]: Unnamed: 0          0
country              US
description    This tremendous 100% varietal wine hails from ...
designation              Martha's Vineyard
points              96
price              235.0
province          California
region_1          Napa Valley
region_2              Napa
variety          Cabernet Sauvignon
winery              Heitz
Name: 0, dtype: object

```

```
In [28]: reviews.iloc[:, 0]
```

```
Out[28]: 0      0
          1      1
          2      2
          3      3
          4      4
          ...
150925    150925
150926    150926
150927    150927
150928    150928
150929    150929
Name: Unnamed: 0, Length: 150930, dtype: int64
```

```
In [29]: reviews.iloc[:3, 0]
```

```
Out[29]: 0      0
          1      1
          2      2
Name: Unnamed: 0, dtype: int64
```

```
In [30]: reviews.iloc[1:3, 0]
```

```
Out[30]: 1      1
          2      2
Name: Unnamed: 0, dtype: int64
```

```
In [31]: reviews.iloc[[0, 1, 2], 0]
```

```
Out[31]: 0      0
          1      1
          2      2
Name: Unnamed: 0, dtype: int64
```

**Finally, it's worth knowing that negative numbers can be used in selection. This will start counting forwards from the end of the values. So for example here are the last five elements of the dataset.**

```
In [32]: reviews.iloc[-5:]
```

```
Out[32]:
```

	Unnamed: 0	country	description	designation	points	price	province	region_1	region_2
150925	150925	Italy	Many people feel Fiano represents southern Ita...	NaN	91	20.0	Southern Italy	Fiano di Avellino	NaN
150926	150926	France	Offers an intriguing nose with ginger, lime an...	Cuvée Prestige	91	27.0	Champagne	Champagne	NaN

	Unnamed: 0	country	description	designation	points	price	province	region_1	region_2
150927	150927	Italy	This classic example comes from a cru vineyard...	Terre di Dora	91	20.0	Southern Italy	Fiano di Avellino	NaN
150928	150928	France	A perfect salmon shade, with scents of peaches...	Grand Brut Rosé	90	52.0	Champagne	Champagne	NaN
150929	150929	Italy	More Pinot Grigios should taste like this. A r...	NaN	90	15.0	Northeastern Italy	Alto Adige	NaN

## Label-based selection

he second paradigm for attribute selection is the one followed by the `loc` operator: label-based selection. In this paradigm, it's the data index value, not its position, which matters.

```
In [33]: reviews.loc[0, 'country']
```

```
Out[33]: 'US'
```

`loc` is conceptually simpler than `loc` because it ignores the dataset's indices. When we use `iloc` we treat the dataset like a big matrix (a list of lists), one that we have to index into by position. `loc`, by contrast, uses the information in the indices to do its work. Since your dataset usually has meaningful indices, it's usually easier to do things using `loc` instead. For example, here's one operation that's much easier using `loc`:

```
In [36]: reviews.loc[:, ['country', 'description', 'points']]
```

```
Out[36]:
```

	country	description	points
0	US	This tremendous 100% varietal wine hails from ...	96
1	Spain	Ripe aromas of fig, blackberry and cassis are ...	96
2	US	Mac Watson honors the memory of a wine once ma...	96
3	US	This spent 20 months in 30% new French oak, an...	96
4	France	This is the top wine from La Bégude, named aft...	95
...	...	...	...
150925	Italy	Many people feel Fiano represents southern Ita...	91





	Unnamed: 0	country	description	designation	points	price	region_1	region_2	
province									
Southern Italy	150925	Italy	Many people feel Fiano represents southern Ita...	NaN	91	20.0	Fiano di Avellino	NaN	
Champagne	150926	France	Offers an intriguing nose with ginger, lime an...	Cuvée Prestige	91	27.0	Champagne	NaN	Char
Southern Italy	150927	Italy	This classic example comes from a cru vineyard...	Terre di Dora	91	20.0	Fiano di Avellino	NaN	
Champagne	150928	France	A perfect salmon shade, with scents of peaches...	Grand Brut Rosé	90	52.0	Champagne	NaN	Char
Northeastern Italy	150929	Italy	More Pinot Grigios should taste like this. A r...	NaN	90	15.0	Alto Adige	NaN	Pino

150930 rows × 10 columns



## Conditional selection

So far we've been indexing various strides of data, using structural properties of the **DataFrame** itself. To do interesting things with the data, however, we often need to ask questions based on conditions.

```
In [39]: reviews.country == 'Italy'
```

```
Out[39]: 0      False
1      False
2      False
3      False
4      False
...
150925  True
150926  False
150927  True
150928  False
```

150929      True  
Name: country, Length: 150930, dtype: bool

```
In [40]: reviews.loc[reviews.country == 'Italy']
```

Out[40]:

	Unnamed: 0	country	description	designation	points	price	province	region_1	regio
10	10	Italy	Elegance, complexity and structure come together...	Ronco della Chiesa	95	80.0	Northeastern Italy	Collio	I
32	32	Italy	Underbrush, scorched earth, menthol and plum s...	Vigna Piaggia	90	NaN	Tuscany	Brunello di Montalcino	I
35	35	Italy	Forest floor, tilled soil, mature berry and a ...	Riserva	90	135.0	Tuscany	Brunello di Montalcino	I
37	37	Italy	Aromas of forest floor, violet, red berry and ...	NaN	90	29.0	Tuscany	Vino Nobile di Montepulciano	I
38	38	Italy	This has a charming nose that boasts rose, vio...	NaN	90	23.0	Tuscany	Chianti Classico	I
...	...	...	...	...	...	...	...	...	...
150920	150920	Italy	Rich and mature aromas of smoke, earth and her...	Brut Riserva	91	19.0	Northeastern Italy	Trento	I
150922	150922	Italy	Made by 30-ish Roberta Borghese high above Man...	Superiore	91	NaN	Northeastern Italy	Colli Orientali del Friuli	I
150925	150925	Italy	Many people feel Fiano represents southern Ita...	NaN	91	20.0	Southern Italy	Fiano di Avellino	I



	Unnamed: 0	country	description	designation	points	price	province	region_1	region
150920	150920	Italy	Rich and mature aromas of smoke, earth and her...	Brut Riserva	91	19.0	Northeastern Italy	Trento	I
150922	150922	Italy	Made by 30-ish Roberta Borghese high above Man...	Superiore	91	NaN	Northeastern Italy	Colli Orientali del Friuli	I
150925	150925	Italy	Many people feel Fiano represents southern Ita...	NaN	91	20.0	Southern Italy	Fiano di Avellino	I
150927	150927	Italy	This classic example comes from a cru vineyard...	Terre di Dora	91	20.0	Southern Italy	Fiano di Avellino	I
150929	150929	Italy	More Pinot Grigios should taste like this. A r...	NaN	90	15.0	Northeastern Italy	Alto Adige	I

7945 rows × 11 columns



```
In [42]: reviews.loc[(reviews.country == 'Italy') | (reviews.points >= 90)]
```

Out[42]:

	Unnamed: 0	country	description	designation	points	price	province	region_1	region
0	0	US	This tremendous 100% varietal wine hails from ...	Martha's Vineyard	96	235.0	California	Napa Valley	Na
1	1	Spain	Ripe aromas of fig, blackberry and cassis are ...	Carodorum Selección Especial Reserva	96	110.0	Northern Spain	Toro	N

	Unnamed: 0	country	description	designation	points	price	province	region_1	region	
	2	2	US	Mac Watson honors the memory of a wine once made by the family	Special Selected Late Harvest	96	90.0	California	Knights Valley	Sonoma
	3	3	US	This spent 20 months in 30% new French oak, and is a blend of Pinot Noir and Chardonnay	Reserve	96	65.0	Oregon	Willamette Valley	Willamette Valley
	4	4	France	This is the top wine from La Bégude, named after the village	La Brûlade	95	66.0	Provence	Bandol	Narbonne
	...	...	...	...	...	...	...	...	...	
150925	150925	Italy	Many people feel Fiano represents southern Italian wine	NaN	91	20.0	Southern Italy	Fiano di Avellino	Naples	
150926	150926	France	Offers an intriguing nose with ginger, lime and a touch of orange	Cuvée Prestige	91	27.0	Champagne	Champagne	Naples	
150927	150927	Italy	This classic example comes from a cru vineyard in the heart of the Fiano di Avellino	Terre di Dora	91	20.0	Southern Italy	Fiano di Avellino	Naples	
150928	150928	France	A perfect salmon shade, with scents of peaches and apricots	Grand Brut Rosé	90	52.0	Champagne	Champagne	Naples	
150929	150929	Italy	More Pinot Grigios should taste like this. A rare find	NaN	90	15.0	Northeastern Italy	Alto Adige	Naples	

63743 rows × 11 columns



```
In [43]: reviews.loc[reviews.price.notnull()]
```

Out[43]:

	Unnamed: 0	country	description	designation	points	price	province	region_1	region
0	0	US	This tremendous 100% varietal wine hails from ...	Martha's Vineyard	96	235.0	California	Napa Valley	Napa
1	1	Spain	Ripe aromas of fig, blackberry and cassis are ...	Carodorum Selección Especial Reserva	96	110.0	Northern Spain	Toro	Napa
2	2	US	Mac Watson honors the memory of a wine once ma...	Special Selected Late Harvest	96	90.0	California	Knights Valley	Sonoma
3	3	US	This spent 20 months in 30% new French oak, an...	Reserve	96	65.0	Oregon	Willamette Valley	Willamette Vall
4	4	France	This is the top wine from La Bégude, named aft...	La Brûlade	95	66.0	Provence	Bandol	Napa
...	...	...	...	...	...	...	...	...	...
150925	150925	Italy	Many people feel Fiano represents southern Ita...	NaN	91	20.0	Southern Italy	Fiano di Avellino	Napa
150926	150926	France	Offers an intriguing nose with ginger, lime an...	Cuvée Prestige	91	27.0	Champagne	Champagne	Napa
150927	150927	Italy	This classic example comes from a cru vineyard...	Terre di Dora	91	20.0	Southern Italy	Fiano di Avellino	Napa
150928	150928	France	A perfect salmon shade, with scents of peaches...	Grand Brut Rosé	90	52.0	Champagne	Champagne	Napa

	Unnamed: 0	country	description	designation	points	price	province	region_1	region
150929	150929	Italy	More Pinot Grigios should taste like this. A r...	NaN	90	15.0	Northeastern Italy	Alto Adige	NaN

137235 rows × 11 columns



with more and more Conditions, thw size of dateset will change to select only data that matches the Conditions

## Assigning data

assigning data to a DataFrame is easy. You can assign either a constant value:

```
In [44]: reviews['critic'] = 'everyone'
reviews['critic']
```

```
Out[44]: 0      everyone
1      everyone
2      everyone
3      everyone
4      everyone
...
150925  everyone
150926  everyone
150927  everyone
150928  everyone
150929  everyone
Name: critic, Length: 150930, dtype: object
```

```
In [46]: reviews['index_backwards'] = range(len(reviews), 0, -1)
reviews['index_backwards']
```

```
Out[46]: 0      150930
1      150929
2      150928
3      150927
4      150926
...
150925      5
150926      4
150927      3
150928      2
150929      1
Name: index_backwards, Length: 150930, dtype: int32
```

## Summary Functions and Maps

the data does not always come out of memory in the format we want it in right out of the bat. Sometimes we have to do some more work ourselves to reformat it for the task at hand. This tutorial will cover different operations we can apply to our data to get the input "just right".

## Summary

This method generates a high-level summary of the attributes of the given column. It is type-aware, meaning that its output changes based on the data type of the input. The output above only makes sense for numerical data; for string data here's what we get:

```
In [6]: reviews.describe()
```

```
Out[6]:
```

	Unnamed: 0	points	price
count	150930.000000	150930.000000	137235.000000
mean	75464.500000	87.888418	33.131482
std	43569.882402	3.222392	36.322536
min	0.000000	80.000000	4.000000
25%	37732.250000	86.000000	16.000000
50%	75464.500000	88.000000	24.000000
75%	113196.750000	90.000000	40.000000
max	150929.000000	100.000000	2300.000000

```
In [9]: reviews.points.mean()
```

```
Out[9]: 87.8884184721394
```

To see a list of unique values we can use the `unique()` function:

```
In [11]: reviews.country.unique()
```

```
Out[11]: array(['US', 'Spain', 'France', 'Italy', 'New Zealand', 'Bulgaria',  
                'Argentina', 'Australia', 'Portugal', 'Israel', 'South Africa',  
                'Greece', 'Chile', 'Morocco', 'Romania', 'Germany', 'Canada',  
                'Moldova', 'Hungary', 'Austria', 'Croatia', 'Slovenia', nan,  
                'India', 'Turkey', 'Macedonia', 'Lebanon', 'Serbia', 'Uruguay',  
                'Switzerland', 'Albania', 'Bosnia and Herzegovina', 'Brazil',  
                'Cyprus', 'Lithuania', 'Japan', 'China', 'South Korea', 'Ukraine',  
                'England', 'Mexico', 'Georgia', 'Montenegro', 'Luxembourg',  
                'Slovakia', 'Czech Republic', 'Egypt', 'Tunisia', 'US-France'],  
              dtype=object)
```

To see a list of unique values and how often they occur in the dataset, we can use the `value_counts()` method:

```
In [12]:
```



```
reviews.country.value_counts()
```

```
Out[12]: US                62397
Italy                23478
France              21098
Spain               8268
Chile               5816
Argentina           5631
Portugal            5322
Australia           4957
New Zealand         3320
Austria             3057
Germany             2452
South Africa        2258
Greece              884
Israel              630
Hungary             231
Canada              196
Romania             139
Slovenia            94
Uruguay             92
Croatia             89
Bulgaria            77
Moldova             71
Mexico              63
Turkey              52
Georgia             43
Lebanon             37
Cyprus              31
Brazil              25
Macedonia           16
Serbia              14
Morocco             12
England             9
Luxembourg          9
Lithuania           8
India               8
Czech Republic      6
Ukraine             5
Switzerland         4
South Korea         4
Bosnia and Herzegovina 4
Slovakia            3
Egypt               3
China               3
Albania             2
Tunisia             2
Montenegro          2
Japan               2
US-France           1
Name: country, dtype: int64
```

## Maps

***A map is a term, borrowed from mathematics, for a function that takes one set of values and "maps" them to another set of values. In data science we often have a need for creating new representations from existing data, or for transforming data from the format it is in now to the format that we want it to be in later. Maps are what handle this work, making them extremely important for getting your work done!***

The function you pass to `map()` should expect a single value from the Series (a point value, in the above example), and return a transformed version of that value. `map()` returns a new Series where all the values have been transformed by your function.

```
In [13]: mean = reviews.points.mean()
reviews.points.map(lambda p: p - mean)
```

```
Out[13]: 0      8.111582
1      8.111582
2      8.111582
3      8.111582
4      7.111582
...
150925  3.111582
150926  3.111582
150927  3.111582
150928  2.111582
150929  2.111582
Name: points, Length: 150930, dtype: float64
```

`apply()` is the equivalent method if we want to transform a whole DataFrame by calling a custom method on each row.

```
In [15]: def remean_points(row):
row.points = row.points - mean
return row

reviews.apply(remean_points, axis='columns')
```

Unnamed: 0		country	description	designation	points	price	province	region_1	region_2
0	0	US	This tremendous 100% varietal wine hails from ...	Martha's Vineyard	8.111582	235.0	California	Napa Valley	Napa
1	1	Spain	Ripe aromas of fig, blackberry and cassis are ...	Carodorum Selección Especial Reserva	8.111582	110.0	Northern Spain	Toro	
2	2	US	Mac Watson honors the memory of a wine once ma...	Special Selected Late Harvest	8.111582	90.0	California	Knights Valley	Sonoma

	Unnamed: 0	country	description	designation	points	price	province	region_1	region_2
3	3	US	This spent 20 months in 30% new French oak, an...	Reserve	8.111582	65.0	Oregon	Willamette Valley	Willamette Valley
4	4	France	This is the top wine from La Bégude, named aft...	La Brûlade	7.111582	66.0	Provence	Bandol	
...	...	...	...	...	...	...	...	...	...
150925	150925	Italy	Many people feel Fiano represents southern Ita...	NaN	3.111582	20.0	Southern Italy	Fiano di Avellino	
150926	150926	France	Offers an intriguing nose with ginger, lime an...	Cuvée Prestige	3.111582	27.0	Champagne	Champagne	
150927	150927	Italy	This classic example comes from a cru vineyard...	Terre di Dora	3.111582	20.0	Southern Italy	Fiano di Avellino	
150928	150928	France	A perfect salmon shade, with scents of peaches...	Grand Brut Rosé	2.111582	52.0	Champagne	Champagne	
150929	150929	Italy	More Pinot Grigios should taste like this. A r...	NaN	2.111582	15.0	Northeastern Italy	Alto Adige	

150930 rows × 11 columns



**Note that `map()` and `apply()` return new, transformed Series and DataFrames, respectively. They don't modify the original data they're called on. If we look at the first row of reviews, we can see that it still has its original points value.**

```
In [16]: reviews.head(1)
```

Out[16]:

Unnamed: 0	country	description	designation	points	price	province	region_1	region_2	variety
0	0	US	This tremendous 100% varietal wine hails from ...	Martha's Vineyard	96	235.0	California	Napa Valley	Napa Cabernet Sauvignon

## Grouping and Sorting

**we want to group our data, and then do something specific to the group the data is in.**

*groupby()* operation. We'll also cover some additional topics, such as more complex ways to index your DataFrames, along with how to sort your data.

```
In [17]: reviews.groupby('points').points.count()
```

```
Out[17]: points
80      898
81     1502
82     4041
83     6048
84    10708
85    12411
86    15573
87    20747
88    17871
89    12921
90    15973
91    10536
92     9241
93     6017
94     3462
95     1716
96      695
97      365
98      131
99       50
100      24
Name: points, dtype: int64
```

```
In [18]: reviews.groupby('points').price.min()
```

```
Out[18]: points
80      5.0
81      5.0
82      5.0
83      4.0
84      4.0
85      4.0
86      4.0
87      6.0
```

```
88      6.0
89      7.0
90      5.0
91      8.0
92     11.0
93     12.0
94     15.0
95     20.0
96     20.0
97     42.0
98     50.0
99     65.0
100     65.0
Name: price, dtype: float64
```

=> *you can also group by more than one column. For an example, here's how we would pick out the best wine by country and province*

```
In [19]: reviews.groupby(['country', 'province']).apply(lambda df: df.loc[df.points.idxmax()])
```

Out[19]:

		Unnamed: 0	country	description	designation	points	price	province	reg
country	province								
Albania	Mirditë	4642	Albania	This garnet-colored wine made from 100% Kallme...	NaN	88	20.0	Mirditë	
Argentina	Mendoza Province	65331	Argentina	If the color doesn't tell the full story, the ...	Nicasia Vineyard	97	120.0	Mendoza Province	Me
	Other	10619	Argentina	Take note, this could be the best wine Colomé ...	Reserva	95	90.0	Other	
Australia	Australia Other	68251	Australia	This big wine presents a sophisticated bouquet...	Yattarna	92	65.0	Australia Other	E Au
	New South Wales	54205	Australia	This wine's deep brassy color suggests honey, ...	Noble One Botrytis	93	32.0	New South Wales	
...	...	...	...	...	...	...	...	...	
Uruguay	Juanico	3160	Uruguay	This mature Bordeaux-style blend is earthy on ...	Preludio Barrel Select Lote N 77	90	45.0	Juanico	

		Unnamed: 0	country	description	designation	points	price	province	reg
country	province								
	Montevideo	3164	Uruguay	Bouza ranks as one of Uruguay's top wineries, ...	Monte Vide Eu Tannat-Merlot-Tempranillo	90	57.0	Montevideo	
	Progreso	6541	Uruguay	Blackberry and plum aromas come with wood-smok...	RPF	89	23.0	Progreso	
	San Jose	70157	Uruguay	While this ranks as one of the best Uruguayan ...	El Preciado Premier Gran Reserva	89	60.0	San Jose	
	Uruguay	132482	Uruguay	They call it Special Barrel, and one sniff tel...	Special Barrel	89	50.0	Uruguay	

455 rows × 11 columns



*agg()*, which lets you run a bunch of different functions on your DataFrame simultaneously. For example, we can generate a simple statistical summary of the dataset as follows:

```
In [24]: reviews.groupby(['country']).price.agg([len, min, max]).head(10)
```

Out[24]:

	len	min	max
country			
Albania	2.0	20.0	20.0
Argentina	5631.0	4.0	250.0
Australia	4957.0	5.0	850.0
Austria	3057.0	8.0	1100.0
Bosnia and Herzegovina	4.0	12.0	13.0
Brazil	25.0	11.0	35.0
Bulgaria	77.0	7.0	28.0
Canada	196.0	12.0	145.0
Chile	5816.0	5.0	400.0
China	3.0	7.0	27.0

## Multi-indexes

A multi-index differs from a regular index in that it has multiple levels. For example:

```
In [9]: countries_reviewed = reviews.groupby(['country', 'province']).description.agg([len])
countries_reviewed
```

```
Out[9]:
```

	country	province	len
	Albania	Mirditë	2
	Argentina	Mendoza Province	4742
		Other	889
	Australia	Australia Other	553
		New South Wales	246
	...	...	...
	Uruguay	Juanico	19
		Montevideo	3
		Progreso	5
		San Jose	15
		Uruguay	18

455 rows × 1 columns

```
In [27]: type(reviews.index) ## RangeIndex
```

```
Out[27]: pandas.core.indexes.range.RangeIndex
```

```
In [26]: mi = countries_reviewed.index ## MultiIndex
type(mi)
```

```
Out[26]: pandas.core.indexes.multi.MultiIndex
```

### for converting back to a regular index

```
In [30]: countries_reviewed.reset_index()
```

```
Out[30]:
```

	country	province	len
0	Albania	Mirditë	2
1	Argentina	Mendoza Province	4742
2	Argentina	Other	889

	country	province	len
3	Australia	Australia Other	553
4	Australia	New South Wales	246
...	...	...	...
450	Uruguay	Juanico	19
451	Uruguay	Montevideo	3
452	Uruguay	Progreso	5
453	Uruguay	San Jose	15
454	Uruguay	Uruguay	18

455 rows × 3 columns

## Sorting

To get data in the order want it in we can sort it ourselves. The `sort_values()` method is handy for this.

```
In [10]: countries_reviewed = countries_reviewed.reset_index()
countries_reviewed.sort_values(by='len')
```

```
Out[10]:
```

	country	province	len
154	Greece	Central Greece	1
207	Greece	Zitsa	1
115	Cyprus	Pafos	1
362	Slovenia	Slovenska Istra	1
213	Hungary	Pannon	1
...	...	...	...
407	Spain	Northern Spain	4892
122	France	Bordeaux	6111
242	Italy	Tuscany	7281
442	US	Washington	9750
422	US	California	44508

455 rows × 3 columns

**To sort by index values, use the companion method `sort_index()`. This method has the same arguments and default order:**

```
In [33]: countries_reviewed.sort_index()
```



```
Out[33]:
```

	country	province	len
0	Albania	Mirditë	2
1	Argentina	Mendoza Province	4742
2	Argentina	Other	889
3	Australia	Australia Other	553
4	Australia	New South Wales	246
...	...	...	...
450	Uruguay	Juanico	19
451	Uruguay	Montevideo	3
452	Uruguay	Progreso	5
453	Uruguay	San Jose	15
454	Uruguay	Uruguay	18

455 rows × 3 columns

**know that you can sort by more than one column at a time:**

```
In [34]: countries_reviewed.sort_values(by=['country', 'len'])
```

```
Out[34]:
```

	country	province	len
0	Albania	Mirditë	2
2	Argentina	Other	889
1	Argentina	Mendoza Province	4742
5	Australia	Queensland	3
7	Australia	Tasmania	47
...	...	...	...
448	Uruguay	Colonia	6
453	Uruguay	San Jose	15
454	Uruguay	Uruguay	18
447	Uruguay	Canelones	19
450	Uruguay	Juanico	19

455 rows × 3 columns

## Dtypes

The data type for a column in a DataFrame or a Series is known as the dtype

Data types tell us something about how pandas is storing the data internally. float64 means that it's using a 64-bit floating point number; int64 means a similarly sized integer instead, and so on.

```
In [5]: reviews.dtypes
```

```
Out[5]: Unnamed: 0      int64
country      object
description  object
designation  object
points       int64
price        float64
province     object
region_1     object
region_2     object
variety      object
winery       object
dtype: object
```

```
In [4]: reviews.price.dtype
```

```
Out[4]: dtype('float64')
```

***One peculiarity to keep in mind (and on display very clearly here) is that columns consisting entirely of strings do not get their own type; they are instead given the object type.***

```
In [7]: reviews.country.dtype
```

```
Out[7]: dtype('O')
```

## Missing data

***Entries missing values are given the value NaN, short for "Not a Number". For technical reasons these NaN values are always of the float64 dtype.***

```
In [5]: reviews[pd.isnull(reviews.country)]
```

```
Out[5]:
```

	Unnamed: 0	country	description	designation	points	price	province	region_1	region_2	va
1133	1133	NaN	Delicate white flowers and a spin of lemon pee...	Askitikos	90	17.0	NaN	NaN	NaN	Assy
1440	1440	NaN	A blend of 60% Syrah, 30% Cabernet Sauvignon a...	Shah	90	30.0	NaN	NaN	NaN	E

	Unnamed: 0	country	description	designation	points	price	province	region_1	region_2	va
68226	68226	NaN	From first sniff to last, the nose never makes...	Piedra Feliz	81	15.0	NaN	NaN	NaN	I
113016	113016	NaN	From first sniff to last, the nose never makes...	Piedra Feliz	81	15.0	NaN	NaN	NaN	I
135696	135696	NaN	From first sniff to last, the nose never makes...	Piedra Feliz	81	15.0	NaN	NaN	NaN	I



In [6]:

```
reviews[pd.notnull(reviews.country)]
```

Out[6]:

	Unnamed: 0	country	description	designation	points	price	province	region_1	region
0	0	US	This tremendous 100% varietal wine hails from ...	Martha's Vineyard	96	235.0	California	Napa Valley	Na
1	1	Spain	Ripe aromas of fig, blackberry and cassis are ...	Carodorum Selección Especial Reserva	96	110.0	Northern Spain	Toro	Na
2	2	US	Mac Watson honors the memory of a wine once ma...	Special Selected Late Harvest	96	90.0	California	Knights Valley	Sonor
3	3	US	This spent 20 months in 30% new French oak, an...	Reserve	96	65.0	Oregon	Willamette Valley	Willamet Vall
4	4	France	This is the top wine from La Bégude, named aft...	La Brûlade	95	66.0	Provence	Bandol	Na

	Unnamed: 0	country	description	designation	points	price	province	region_1	region
	...	...	...	...	...	...	...	...	...
150925	150925	Italy	Many people feel Fiano represents southern Ita...	NaN	91	20.0	Southern Italy	Fiano di Avellino	NaN
150926	150926	France	Offers an intriguing nose with ginger, lime an...	Cuvée Prestige	91	27.0	Champagne	Champagne	NaN
150927	150927	Italy	This classic example comes from a cru vineyard...	Terre di Dora	91	20.0	Southern Italy	Fiano di Avellino	NaN
150928	150928	France	A perfect salmon shade, with scents of peaches...	Grand Brut Rosé	90	52.0	Champagne	Champagne	NaN
150929	150929	Italy	More Pinot Grigios should taste like this. A r...	NaN	90	15.0	Northeastern Italy	Alto Adige	NaN

150925 rows × 11 columns



```
In [9]: reviews.isnull().sum()
```

```
Out[9]: Unnamed: 0      0
country      5
description  0
designation  45735
points      0
price      13695
province     5
region_1    25060
region_2    89977
variety      0
winery       0
dtype: int64
```

```
In [10]: reviews.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150930 entries, 0 to 150929
Data columns (total 11 columns):
```

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	150930 non-null	int64
1	country	150925 non-null	object
2	description	150930 non-null	object
3	designation	105195 non-null	object
4	points	150930 non-null	int64
5	price	137235 non-null	float64
6	province	150925 non-null	object
7	region_1	125870 non-null	object
8	region_2	60953 non-null	object
9	variety	150930 non-null	object
10	winery	150930 non-null	object

dtypes: float64(1), int64(2), object(8)  
memory usage: 12.7+ MB

**Replacing missing values is a common operation. Pandas provides a really handy method for this problem: `fillna()`. `fillna()` provides a few different strategies for mitigating such data. For example, we can simply replace each NaN with an "Unknown":**

```
In [11]: reviews.region_2.fillna("Unknown")
```

```
Out[11]: 0          Napa
1          Unknown
2          Sonoma
3    Willamette Valley
4          Unknown
...
150925          Unknown
150926          Unknown
150927          Unknown
150928          Unknown
150929          Unknown
Name: region_2, Length: 150930, dtype: object
```

**Alternatively, we may have a non-null value that we would like to replace. For example, suppose that since this dataset was published, reviewer Kerin O'Keefe has changed her country from US to United States. One way to reflect this in the dataset is using the `replace()` method:**

```
In [15]: reviews.country.replace("US", "United States")
```

```
Out[15]: 0          United States
1             Spain
2          United States
3          United States
4             France
...
150925          Italy
150926          France
150927          Italy
150928          France
150929          Italy
Name: country, Length: 150930, dtype: object
```

**There are many ways to deal with NULL values. But this course was satisfied with these methods above**

# Renaming and Combining

Oftentimes data will come to us with column names, index names, or other naming conventions that we are not satisfied with. In that case, you'll learn how to use pandas functions to change the names of the offending entries to something better.

You'll also explore how to combine data from multiple DataFrames and/or Series.

## Renaming

The first function we'll introduce here is rename(), which lets you change index names and/or column names. For example, to change the points column in our dataset to score, we would do:

```
In [17]: reviews.columns
```

```
Out[17]: Index(['Unnamed: 0', 'country', 'description', 'designation', 'points', 'price', 'province', 'region_1', 'region_2', 'variety'], dtype='object')
```

```
In [24]: reviews.drop('Unnamed: 0', inplace=True, axis=1)
```

```
In [28]: reviews.rename(columns={'points': 'score'})
```

Out[28]:	country	description	designation	score	price	province	region_1	region_2	varie
0	US	This tremendous 100% varietal wine hails from ...	Martha's Vineyard	96	235.0	California	Napa Valley	Napa	Cabern Sauvign
1	Spain	Ripe aromas of fig, blackberry and cassis are ...	Carodorum Selección Especial Reserva	96	110.0	Northern Spain	Toro	NaN	Tinta c To
2	US	Mac Watson honors the memory of a wine once ma...	Special Selected Late Harvest	96	90.0	California	Knights Valley	Sonoma	Sauvign Bla
3	US	This spent 20 months in 30% new French oak, an...	Reserve	96	65.0	Oregon	Willamette Valley	Willamette Valley	Pinot Nc



fields	Unnamed: 0	country	description	designation	points	price	province	region_1	region
wines									
0	0	US	This tremendous 100% varietal wine hails from ...	Martha's Vineyard	96	235.0	California	Napa Valley	Napa
1	1	Spain	Ripe aromas of fig, blackberry and cassis are ...	Carodorum Selección Especial Reserva	96	110.0	Northern Spain	Toro	Napa
2	2	US	Mac Watson honors the memory of a wine once made ...	Special Selected Late Harvest	96	90.0	California	Knights Valley	Sonoma
3	3	US	This spent 20 months in 30% new French oak, and ...	Reserve	96	65.0	Oregon	Willamette Valley	Willamette Valley
4	4	France	This is the top wine from La Bégude, named after ...	La Brûlade	95	66.0	Provence	Bandol	Napa
...	...	...	...	...	...	...	...	...	...
150925	150925	Italy	Many people feel Fiano represents southern Italian ...	NaN	91	20.0	Southern Italy	Fiano di Avellino	Napa
150926	150926	France	Offers an intriguing nose with ginger, lime and ...	Cuvée Prestige	91	27.0	Champagne	Champagne	Napa
150927	150927	Italy	This classic example comes from a cru vineyard ...	Terre di Dora	91	20.0	Southern Italy	Fiano di Avellino	Napa





	country	description	designation	points	price	province	region_1	region_2	vari
150925	Italy	Many people feel Fiano represents southern Ita...	NaN	91	20.0	Southern Italy	Fiano di Avellino	NaN	WI Bl
150926	France	Offers an intriguing nose with ginger, lime an...	Cuvée Prestige	91	27.0	Champagne	Champagne	NaN	Champa Bl
150927	Italy	This classic example comes from a cru vineyard...	Terre di Dora	91	20.0	Southern Italy	Fiano di Avellino	NaN	WI Bl
150928	France	A perfect salmon shade, with scents of peaches...	Grand Brut Rosé	90	52.0	Champagne	Champagne	NaN	Champa Bl
150929	Italy	More Pinot Grigios should taste like this. A r...	NaN	90	15.0	Northeastern Italy	Alto Adige	NaN	Pinot Gri

150930 rows × 10 columns



## Combining

**When performing operations on a dataset, we will sometimes need to combine different DataFrames and/or Series in non-trivial ways. Pandas has three core methods for doing this. In order of increasing complexity, these are *concat()*, *join()*, and *merge()*.**

**The simplest combining method is *concat()*. Given a list of elements, this function will smush those elements together along an axis. This is useful when we have data in different DataFrame or Series objects but having the same fields (columns).**

```
In [37]: df1 = pd.DataFrame({'Yes': [10,20,30,40,50], 'No': [60,70,80,90,100]})
```

```
In [38]: df2 = pd.DataFrame({'Yes': [60,70,80,90,100], 'No': [10,20,30,40,50]})
```

```
In [39]: pd.concat([df1, df2])
```

Out[39]:

	Yes	No
<b>0</b>	10	60
<b>1</b>	20	70
<b>2</b>	30	80
<b>3</b>	40	90
<b>4</b>	50	100
<b>0</b>	60	10
<b>1</b>	70	20
<b>2</b>	80	30
<b>3</b>	90	40
<b>4</b>	100	50

```
In [45]: df1['maybe'] = [15,25,35,45,55]
df1['sure'] = [100,100,100,100,100]
```

```
In [46]: df1
```

```
Out[46]:
```

	Yes	No	maybe	sure
<b>0</b>	10	60	15	100
<b>1</b>	20	70	25	100
<b>2</b>	30	80	35	100
<b>3</b>	40	90	45	100
<b>4</b>	50	100	55	100

```
In [47]: df2['maybe'] = [65,75,85,95,105]
df2['sure'] = [100,100,100,100,100]
```

```
In [48]: df2
```

```
Out[48]:
```

	Yes	No	maybe	sure
<b>0</b>	60	10	65	100
<b>1</b>	70	20	75	100
<b>2</b>	80	30	85	100
<b>3</b>	90	40	95	100
<b>4</b>	100	50	105	100

The middlemost combiner in terms of complexity is `join()`. `join()` lets you combine different `DataFrame` objects which have an index in common.

```
In [56]: left = df1.set_index(['Yes'])
right = df2.set_index(['Yes'])

left.join(left, lsuffix='df1', rsuffix='df2')
```

```
Out[56]:
```

	<b>Nodf1</b>	<b>maybedf1</b>	<b>suredf1</b>	<b>Nodf2</b>	<b>maybedf2</b>	<b>suredf2</b>
<b>Yes</b>						
<b>10</b>	60	15	100	60	15	100
<b>20</b>	70	25	100	70	25	100
<b>30</b>	80	35	100	80	35	100
<b>40</b>	90	45	100	90	45	100
<b>50</b>	100	55	100	100	55	100

# Congratulations!